

FML_Hierarchical Clustering

Yeswanth Siripurapu

2023-11-29

Load Data Set and Libraries #First, we will load all of the packages needed for this task. "ISLR", "caret", "dplyr", "tidyverse", "factoextra", "ggplot2", "proxy", "NbClust", "ppclust", "dendextend", "cluster", and "tinytex" will be loaded for this problem.

```
library(cluster)
library(ISLR)
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Warning: package 'ggplot2' was built under R version 4.3.2
```

```
## Loading required package: lattice
```

```
## Warning: package 'lattice' was built under R version 4.3.2
```

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(tidyverse)
```

```
## — Attaching core tidyverse packages — tidyverse 2.0.0 —
## ✓ forcats 1.0.0   ✓ stringr 1.5.0
## ✓ lubridate 1.9.3 ✓ tibble 3.2.1
## ✓ purrr 1.0.2   ✓ tidyr 1.3.0
## ✓ readr 2.1.4
```

```
## — Conflicts — tidyverse_conflicts() —
## ✖ dplyr::filter() masks stats::filter()
## ✖ dplyr::lag() masks stats::lag()
## ✖ purrr::lift() masks caret::lift()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(factoextra)
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
library(ggplot2)
library(proxy)
```

```
##
## Attaching package: 'proxy'
##
## The following objects are masked from 'package:stats':
##
##   as.dist, dist
##
## The following object is masked from 'package:base':
##
##   as.matrix
```

```
library(NbClust)
library(ppclust)
```

```
## Warning: package 'ppclust' was built under R version 4.3.2
```

```
library(dendextend)
```

```
## Warning: package 'dendextend' was built under R version 4.3.2
```

```
##
## -----
## Welcome to dendextend version 1.17.1
## Type citation('dendextend') for how to cite the package.
##
## Type browseVignettes(package = 'dendextend') for the package vignette.
## The github page is: https://github.com/talgalili/dendextend/
##
## Suggestions and bug-reports can be submitted at: https://github.com/talgalili/dendextend/issues
## You may ask questions at stackoverflow, use the r and dendextend tags:
##   https://stackoverflow.com/questions/tagged/dendextend
##
## To suppress this message use: suppressPackageStartupMessages(library(dendextend))
## -----
##
##
## Attaching package: 'dendextend'
##
## The following object is masked from 'package:stats':
##
##   cutree
```

```
library(tinytex)
```

```
## Warning: package 'tinytex' was built under R version 4.3.2
```

#The “cereal” data set will then be imported into the RStudio environment.

```
# Import data set from BlackBoard into the RStudio environment

cereal <- read.csv("D:/Cereals.csv")
```

Review Data Structure #A summary of the data set will be displayed for inspection.Examine the first few rows of the data set

```
head(cereal)
```

```
##
##           name mfr type calories protein fat sodium fiber carbo
## 1      100%_Bran  N   C       70      4  1   130  10.0   5.0
## 2    100%_Natural_Bran  Q   C      120      3  5    15   2.0   8.0
## 3         All-Bran  K   C       70      4  1   260   9.0   7.0
## 4 All-Bran_with_Extra_Fiber  K   C       50      4  0   140  14.0   8.0
## 5        Almond_Delight  R   C      110      2  2   200   1.0  14.0
## 6  Apple_Cinnamon_Cheerios  G   C      110      2  2   180   1.5  10.5
##  sugars potass vitamins shelf weight cups  rating
## 1      6    280      25    3      1 0.33 68.40297
## 2      8    135       0    3      1 1.00 33.98368
## 3      5    320      25    3      1 0.33 59.42551
## 4      0    330      25    3      1 0.50 93.70491
## 5      8     NA      25    3      1 0.75 34.38484
## 6     10     70      25    1      1 0.75 29.50954
```

```
# Look into the data set's structure.
str(cereal)
```

```
## 'data.frame':    77 obs. of  16 variables:
## $ name      : chr  "100% Bran" "100% Natural Bran" "All-Bran" "All-Bran_with_Extra_Fiber" ...
## $ mfr       : chr  "N" "Q" "K" "K" ...
## $ type      : chr  "C" "C" "C" "C" ...
## $ calories  : int   70 120 70 50 110 110 110 130 90 90 ...
## $ protein   : int    4 3 4 4 2 2 2 3 2 3 ...
## $ fat       : int    1 5 1 0 2 2 0 2 1 0 ...
## $ sodium    : int   130 15 260 140 200 180 125 210 200 210 ...
## $ fiber     : num    10 2 9 14 1 1.5 1 2 4 5 ...
## $ carbo     : num    5 8 7 8 14 10.5 11 18 15 13 ...
## $ sugars    : int    6 8 5 0 8 10 14 8 6 5 ...
## $ potass    : int   280 135 320 330 NA 70 30 100 125 190 ...
## $ vitamins  : int    25 0 25 25 25 25 25 25 25 ...
## $ shelf     : int    3 3 3 3 3 1 2 3 1 3 ...
## $ weight    : num    1 1 1 1 1 1 1 1.33 1 1 ...
## $ cups      : num    0.33 1 0.33 0.5 0.75 0.75 1 0.75 0.67 0.67 ...
## $ rating    : num   68.4 34 59.4 93.7 34.4 ...
```

```
# Investigate the summary of the data set
summary(cereal)
```

```
##      name                mfr                type                calories
## Length:77             Length:77             Length:77             Min.   : 50.0
## Class :character       Class :character       Class :character       1st Qu.:100.0
## Mode  :character       Mode  :character       Mode  :character       Median :110.0
##                                     Mean    :106.9
##                                     3rd Qu.:110.0
##                                     Max.    :160.0
##
##      protein            fat                sodium            fiber
## Min.   :1.000          Min.   :0.000          Min.   : 0.0          Min.   : 0.000
## 1st Qu.:2.000          1st Qu.:0.000          1st Qu.:130.0        1st Qu.: 1.000
## Median :3.000          Median :1.000          Median :180.0        Median : 2.000
## Mean   :2.545          Mean   :1.013          Mean   :159.7        Mean   : 2.152
## 3rd Qu.:3.000          3rd Qu.:2.000          3rd Qu.:210.0        3rd Qu.: 3.000
## Max.   :6.000          Max.   :5.000          Max.   :320.0        Max.   :14.000
##
##      carbo              sugars              potass              vitamins
## Min.   : 5.0           Min.   : 0.000          Min.   : 15.00         Min.   : 0.00
## 1st Qu.:12.0           1st Qu.: 3.000          1st Qu.: 42.50         1st Qu.: 25.00
## Median :14.5           Median : 7.000          Median : 90.00         Median : 25.00
## Mean   :14.8           Mean   : 7.026          Mean   : 98.67         Mean   : 28.25
## 3rd Qu.:17.0           3rd Qu.:11.000         3rd Qu.:120.00        3rd Qu.: 25.00
## Max.   :23.0           Max.   :15.000          Max.   :330.00         Max.   :100.00
## NA's   :1              NA's   :1              NA's   :2
##      shelf              weight              cups              rating
## Min.   :1.000          Min.   :0.50           Min.   :0.250          Min.   :18.04
## 1st Qu.:1.000          1st Qu.:1.00           1st Qu.:0.670          1st Qu.:33.17
## Median :2.000          Median :1.00           Median :0.750          Median :40.40
## Mean   :2.208          Mean   :1.03           Mean   :0.821          Mean   :42.67
## 3rd Qu.:3.000          3rd Qu.:1.00           3rd Qu.:1.000          3rd Qu.:50.83
## Max.   :3.000          Max.   :1.50           Max.   :1.500          Max.   :93.70
##
```

```
#Data Preprocessing
#The data will be scaled prior to removing the NA(Null) values from the data set.
# Create duplicate of data set for preprocessing

cereal_scaled <- cereal
# Scale the data set prior to placing it into a clustering algorithm
cereal_scaled[, c(4:16)] <- scale(cereal[, c(4:16)])
# Remove NA values from data set
cereal_preprocessed <- na.omit(cereal_scaled)

# Review the scaled data set with NA's removed
head(cereal_preprocessed)
```

```
##               name mfr type   calories   protein      fat
## 1          100% Bran    N    C -1.8929836  1.3286071 -0.01290349
## 2      100%_Natural_Bran    Q    C  0.6732089  0.4151897  3.96137277
## 3              All-Bran    K    C -1.8929836  1.3286071 -0.01290349
## 4 All-Bran_with_Extra_Fiber    K    C -2.9194605  1.3286071 -1.00647256
## 6   Apple_Cinnamon_Cheerios    G    C  0.1599704 -0.4982277  0.98066557
## 7       Apple_Jacks        K    C  0.1599704 -0.4982277 -1.00647256
##      sodium      fiber      carbo      sugars      potass      vitamins      shelf
## 1 -0.3539844  3.29284661 -2.5087829 -0.2343906  2.5753685 -0.1453172  0.9515734
## 2 -1.7257708 -0.06375361 -1.7409943  0.2223705  0.5160205 -1.2642598  0.9515734
## 3  1.1967306  2.87327158 -1.9969238 -0.4627711  3.1434645 -0.1453172  0.9515734
## 4 -0.2346986  4.97114672 -1.7409943 -1.6046739  3.2854885 -0.1453172  0.9515734
## 6  0.2424445 -0.27354112 -1.1011705  0.6791317 -0.4071355 -0.1453172 -1.4507595
## 7 -0.4136273 -0.48332864 -0.9732057  1.5926539 -0.9752315 -0.1453172 -0.2495930
##      weight      cups      rating
## 1 -0.1967771 -2.1100340  1.8321876
## 2 -0.1967771  0.7690100 -0.6180571
## 3 -0.1967771 -2.1100340  1.1930986
## 4 -0.1967771 -1.3795303  3.6333849
## 6 -0.1967771 -0.3052601 -0.9365625
## 7 -0.1967771  0.7690100 -0.6756899
```

#After pre-processing and scaling the data, the total number of observations went from 77 to 74. Therefore, there were only 3 records with an "NA" value

Assignment Task A #Apply hierarchical clustering to the data using Euclidean distance to the normalized measurements. Use Agnes to compare the clustering from single linkage, complete linkage, average linkage, and Ward. Choose the best method."

#Single Linkage:

Create the dissimilarity matrix for the numeric values in the data set via Euclidean distance measurements

```
cereal_d_euclidean <- dist(cereal_preprocessed[ , c(4:16)], method ="euclidean")
```

```
# Perform hierarchical clustering via the single linkage method
```

```
ag_hc_single <- agnes(cereal_d_euclidean, method = "single")
```

```
# Plot the results of the different methods
```

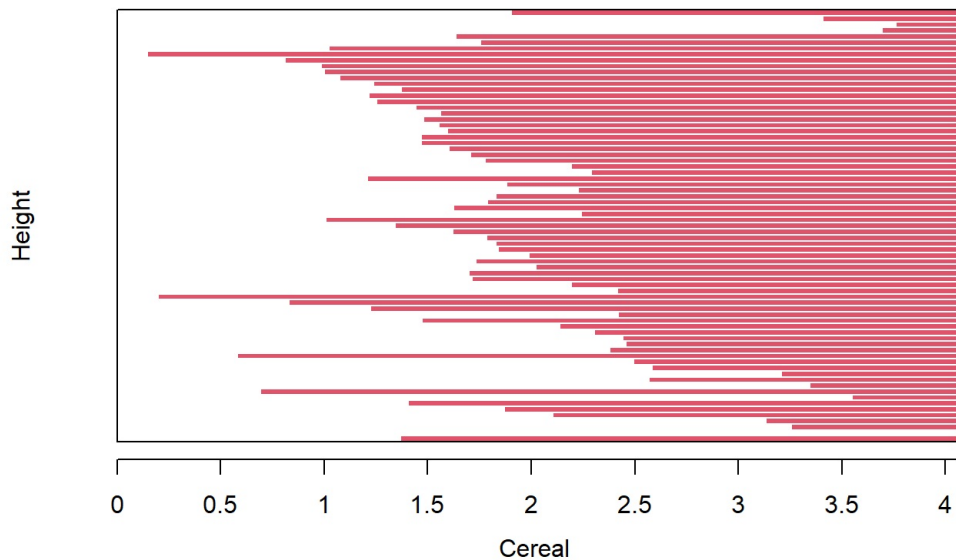
```
plot(ag_hc_single,
     main = "Customer Cereal Ratings - AGNES - Single Linkage Method",
     xlab = "Cereal",
     ylab = "Height",
     cex.axis = 1,
     cex = 0.55,
     hang = -1)
```

```
## Warning in plot.window(xlim, ylim, log = log, ...): "hang" is not a graphical
## parameter
```

```
## Warning in title(main = main, sub = sub, xlab = xlab, ylab = ylab, ...): "hang"
## is not a graphical parameter
```


```
## Warning in axis(1, at = at.vals, labels = lab.vals, ...): "hang" is not a
## graphical parameter
```

Customer Cereal Ratings - AGNES - Single Linkage Method



Agglomerative Coefficient = 0.61

Customer Cereal Ratings - AGNES - Single Linkage Method



Agglomerative Coefficient = 0.61


```
#Complete Linkage:
# Perform hierarchical clustering via the complete linkage method
ag_hc_complete <- agnes(cereal_d_euclidean, method = "complete")
# Plot the results of the different methods
plot(ag_hc_complete,
  main = "Customer Cereal Ratings - AGNES - Complete Linkage Method",
  xlab = "Cereal",
  ylab = "Height",
  cex.axis = 1,
  cex = 0.55,
  hang = -1)
```

```
## Warning in plot.window(xlim, ylim, log = log, ...): "hang" is not a graphical
## parameter
```

```
## Warning in title(main = main, sub = sub, xlab = xlab, ylab = ylab, ...): "hang"
## is not a graphical parameter
```


```
## Warning in axis(1, at = at.vals, labels = lab.vals, ...): "hang" is not a
## graphical parameter
```

Customer Cereal Ratings - AGNES - Complete Linkage Method



Agglomerative Coefficient = 0.84

Customer Cereal Ratings - AGNES - Complete Linkage Method



Agglomerative Coefficient = 0.84


```
#Average Linkage:
# Perform hierarchical clustering via the average linkage method
ag_hc_average <- agnes(cereal_d_euclidean, method = "average")
# Plot the results of the different methods
plot(ag_hc_average,
  main = "Customer Cereal Ratings - AGNES - Average Linkage Method",
  xlab = "Cereal",
  ylab = "Height",
  cex.axis = 1,
  cex = 0.55,
  hang = -1)
```

```
## Warning in plot.window(xlim, ylim, log = log, ...): "hang" is not a graphical
## parameter
```

```
## Warning in title(main = main, sub = sub, xlab = xlab, ylab = ylab, ...): "hang"
## is not a graphical parameter
```


```
## Warning in axis(1, at = at.vals, labels = lab.vals, ...): "hang" is not a
## graphical parameter
```

Customer Cereal Ratings - AGNES - Average Linkage Method



Agglomerative Coefficient = 0.78

Customer Cereal Ratings - AGNES - Average Linkage Method



Agglomerative Coefficient = 0.78


```
#Ward Method:
# Perform hierarchical clustering via the ward linkage method
ag_hc_ward <- agnes(cereal_d_euclidean, method = "ward")
# Plot the results of the different methods
plot(ag_hc_ward,
  main = "Customer Cereal Ratings - AGNES - Ward Linkage Method",
  xlab = "Cereal",
  ylab = "Height",
  cex.axis = 1,
  cex = 0.55,
  hang = -1)
```

```
## Warning in plot.window(xlim, ylim, log = log, ...): "hang" is not a graphical
## parameter
```

```
## Warning in title(main = main, sub = sub, xlab = xlab, ylab = ylab, ...): "hang"
## is not a graphical parameter
```


```
## Warning in axis(1, at = at.vals, labels = lab.vals, ...): "hang" is not a
## graphical parameter
```

Customer Cereal Ratings - AGNES - Ward Linkage Method



Agglomerative Coefficient = 0.9

Customer Cereal Ratings - AGNES - Ward Linkage Method



Agglomerative Coefficient = 0.9

#The best clustering method would be based on the agglomerative coefficient that is returned from each method. The closer the value is to 1.0, the closer the clustering structure is. Therefore, the method with the value closest to 1.0 will be chosen.

#Single Linkage: 0.61 #Complete Linkage: 0.84 #Average Linkage: 0.78 #Ward Method: 0.90
#As a result, the Ward method will be chosen as the best clustering model in this problem.

#Assignment Task B

#“How many clusters would you choose?”

#To determine the appropriate number of clusters, we will use the elbow and silhouette methods.

#Elbow Method:


Determine the optimal number of clusters for the dataset via the Elbow method

```
fviz_nbclust(cereal_preprocessed[, c(4:16)], hcute, method = "wss", k.max = 25) +
```

```
labs(title = "Optimal Number of Clusters - Elbow Method") +
```


```
geom_vline(xintercept = 12, linetype = 2)
```


Optimal Number of Clusters - Elbow Method



```
#Silhouette Method:
# Determine the optimal number of clusters for the dataset via the silhouette method
fviz_nbclust(cereal_preprocessed[, c(4:16)],
             hcut,
             method = "silhouette",
             k.max = 25) +
labs(title = "Optimal Number of Clusters - Silhouette Method")
```

Optimal Number of Clusters - Silhouette Method




```
#Based on the agreement of the silhouette and elbow method, the appropriate number of clusters would be 12 in this case.
#Below we will outline the 12 clusters on the hierarchical tree
# Plot of the Ward hierarchical tree with the 12 clusters outlined for reference
plot(ag_hc_ward,
     main = "AGNES - Ward Linkage Method - 12 Clusters Outlined",
     xlab = "Cereal",
     ylab = "Height",
     cex.axis = 1,
     cex = 0.55,
     hang = -1)
```

```
## Warning in plot.window(xlim, ylim, log = log, ...): "hang" is not a graphical
## parameter
```

```
## Warning in title(main = main, sub = sub, xlab = xlab, ylab = ylab, ...): "hang"
## is not a graphical parameter
```


```
## Warning in axis(1, at = at.vals, labels = lab.vals, ...): "hang" is not a
## graphical parameter
```

AGNES - Ward Linkage Method - 12 Clusters Outlined



Agglomerative Coefficient = 0.9

AGNES - Ward Linkage Method - 12 Clusters Outlined



Agglomerative Coefficient = 0.9

#Assignment Task C

#“Comment on the structure of the clusters and on their stability. Hint: To check stability, partition the data and see how well clusters formed based on one part apply to the other part. To do this:

#1. Cluster partition A #2. Use the cluster centroids from A to assign each record in partition B (each record is assigned to the cluster with the closest centroid).

#3.Assess how consistent the cluster assignments are compared to the assignments based on all the data”

#All Data Assigned Clusters:

#The assigned clusters for all data sets will be in “cereal_preprocessed_1”:

```
# Cut the tree into 12 clusters for analysis
ward_clusters_12 <- cutree(ag_hc_ward, k = 12)
```

```
# Add the assigned cluster to the preprocessed data set
cereal_preprocessed_1 <- cbind(cluster = ward_clusters_12,
cereal_preprocessed)
```

#Partition Data:

#To check stability of clusters, the data set will be split into a 70/30 partition. The 70% will be used to create cluster assignments again, and then the remaining 30% will be assigned based on their closest centroid.

```
# Set the seed for randomized functions
set.seed(982579)
```

```
# Split the data into 70% partition A and 30% partition B
cerealIndex <- createDataPartition(cereal_preprocessed$protein, p=0.3, list =F)
cereal_preprocessed_PartitionB <- cereal_preprocessed[cerealIndex, ]
cereal_preprocessed_PartitionA <- cereal_preprocessed[-cerealIndex,]
```

#Re-Run Clustering with Partitioned Data:

#For the purposes of this task, we will assume the same K value (12) and ward clustering method to determine the stability of the clusters. We will then assign clusters to the nearest points in Partition B (for clusters 1 to 12).

```
# Create the dissimilarity matrix for the numeric values in the partitioned data set via Euclidean distance measurements
cereal_d_euclidean_A <- dist(cereal_preprocessed_PartitionA[ , c(4:16)],
method = "euclidean")
```

```
# Perform hierarchical clustering via the ward linkage method on partitioned data
ag_hc_ward_A <- agnes(cereal_d_euclidean_A, method = "ward")
```


```
# Plot the results of the different methods
plot(ag_hc_ward_A,
      main = "Customer Cereal Ratings - Ward Linkage Method - Partition A",
      xlab = "Cereal",
      ylab = "Height",
      cex.axis = 1,
      cex = 0.55,
      hang = -1)
```

```
## Warning in plot.window(xlim, ylim, log = log, ...): "hang" is not a graphical
## parameter
```

```
## Warning in title(main = main, sub = sub, xlab = xlab, ylab = ylab, ...): "hang"
## is not a graphical parameter
```


```
## Warning in axis(1, at = at.vals, labels = lab.vals, ...): "hang" is not a
## graphical parameter
```

Customer Cereal Ratings - Ward Linkage Method - Partition A



Agglomerative Coefficient = 0.88

Customer Cereal Ratings - Ward Linkage Method - Partition A



Agglomerative Coefficient = 0.88

```

# Cut the tree into 12 clusters for analysis
ward_clusters_12_A <- cutree(ag_hc_ward_A, k = 12)

# Add the assigned cluster to the preprocessed data set
cereal_preprocessed_A <- cbind(cluster = ward_clusters_12_A,
cereal_preprocessed_PartitionA)

#The centroids for each of the clusters will need to be calculated, so we can find the closest centroid for the d
ata points in partition B.

# Find the centroids for the re-ran Ward hierarchical clustering
ward_Centroids_A <- aggregate(cereal_preprocessed_A[ , 5:17],
list(cereal_preprocessed_A$cluster), mean)
ward_Centroids_A <- data.frame(Cluster = ward_Centroids_A[ , 1], Centroid =
rowMeans(ward_Centroids_A[ , -c(1:4)]))
ward_Centroids_A <- ward_Centroids_A$Centroid

# Calculate Centers of Partition B data set
cereal_preprocessed_PartitionB_centers <-
data.frame(cereal_preprocessed_PartitionB[ , 1:3], Center =
rowMeans(cereal_preprocessed_PartitionB[ , 4:16]))

# Calculate the distance between the centers of partition A and the values of partition B
B_to_A_centers <- dist(ward_Centroids_A,
cereal_preprocessed_PartitionB_centers$Center, method = "euclidean")

# Assign the clusters based on the minimum distance to cluster centers
cereal_preprocessed_B <- cbind(cluster =
c(4,8,7,3,5,6,7,11,11,10,8,5,10,1,10,1,4,12,12,7,7,1,4,9),
cereal_preprocessed_PartitionB)

# Combine partitions A and B for comparision to original clusters
cereal_preprocessed_2 <- rbind(cereal_preprocessed_A, cereal_preprocessed_B)
cereal_preprocessed_1 <-
cereal_preprocessed_1[order(cereal_preprocessed_1$name), ]
cereal_preprocessed_2 <-
cereal_preprocessed_2[order(cereal_preprocessed_2$name), ]

#Now that the data has been assigned by both methods (full data and partitioned data), we can compare the number
of matching assignments to see the stability of the clusters.
sum(cereal_preprocessed_1$cluster == cereal_preprocessed_2$cluster)

```

```
## [1] 14
```

#From this result, it can be stated that the clusters are not very stable. With 70% of the data available, the resulting assignments were only identical for 35 out of the 74 observations. This results in a 47% repeatability of assignment.

Visualize the cluster assignments to see any difference between the two

Plot of original hierarchical clustering algorithm

```

ggplot(data = cereal_preprocessed_1, aes(cereal_preprocessed_1$cluster)) +
  geom_bar(fill = "yellow3") +
  labs(title="Count of Cluster Assignments - All Original Data") +
  labs(x="Cluster Assignment", y="Count") +
  guides(fill=FALSE) +
  scale_x_continuous(breaks=c(1:12)) +
  scale_y_continuous(breaks=c(5,10,15,20), limits = c(0,25))


```

```

## Warning: The `<scale>` argument of `guides()` cannot be `FALSE`. Use "none" instead as
## of ggplot2 3.3.4.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.


```

Count of Cluster Assignments - All Original Data



```
# Plot of algorithm that was partitioned prior to assigning the remaining data
ggplot(data = cereal_preprocessed_2, aes(cereal_preprocessed_2$cluster)) +
  geom_bar(fill = "yellow3") +
  labs(title="Count of Cluster Assignments - Partitioned Data") +
  labs(x="Cluster Assignment", y="Count") +
  guides(fill=FALSE) +
  scale_x_continuous(breaks=c(1:12)) +
  scale_y_continuous(breaks=c(5,10,15,20), limits = c(0,25))
```

Count of Cluster Assignments - Partitioned Data



#When we use partitioned data, we can see that Cluster 3 shrinks dramatically. As a result, several of the other clusters grew in size. When the data is partitioned, the clusters appear to be more uniformly spread over the 12 clusters, according to the graphic.

Assignment Task D

#“The elementary public schools would like to choose a set of cereals to include in their daily cafeterias. Every day a different cereal is offered, but all cereals should support a healthy diet. For this goal, you are requested to find a cluster of “healthy cereals.” Should the data be normalized? If not, how should they be used in the cluster analysis?”

#Normalizing the data would be inappropriate in this circumstance. This is due to the fact that the scaling or normalization of cereal nutrition data is dependent on the exact sample of cereal being investigated. As a result, the data set gathered may only include cereals with extremely high sugar content but poor fiber, iron, and other nutritious features. When the data inside the sample set is scaled or normalized, estimating the nutritional value of the cereal for a child becomes impossible. An uneducated observer could believe that a cereal with an iron score of 0.999 provides virtually all of a child’s needed iron; yet, it could just be the best among the worst in the sample set, providing little to no iron.

#As a result, a better way to preprocess the data would be to make it a ratio to a child's daily prescribed calories, fiber, carbohydrates, and so on. This would allow analysts to make more educated decisions on clusters when reviewing them, while not allowing a few larger variables to trump distance calculations. An analyst might evaluate the cluster averages to calculate what percentage of a student's daily needed nutrition would come from XX cereal while evaluating the clusters. This would enable the staff to make more educated selections about which "healthy" cereal clusters to select.